

A quick look into data augmentations in nnUNetv2

Copyright information

License: CC-BY-NC-ND

0. Content

1. Introduction
 - a. General ideas
 - b. Where it is called
2. Key components
 - a. Rotation policies (adaptive)
 - b. Spatial transforms (almost fixed)
 - c. Intensity (photometric) transforms (almost fixed)
 - d. Additional steps for cascaded training (low-res + full_res)
3. Comments
4. Beyond pre-configured data augmentations, taking CMRxMotion and FeTA-2022 as examples

1. Introduction

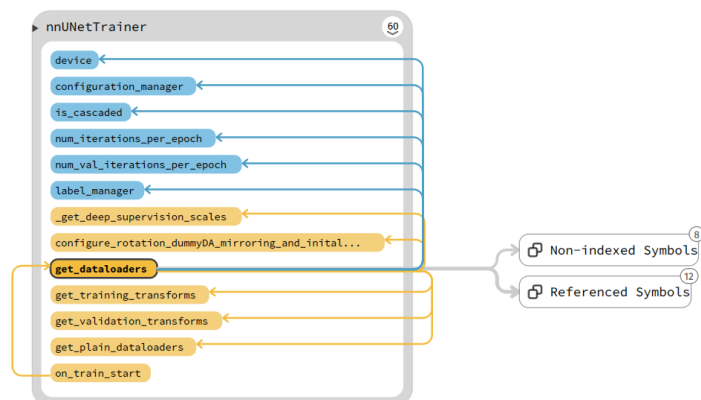
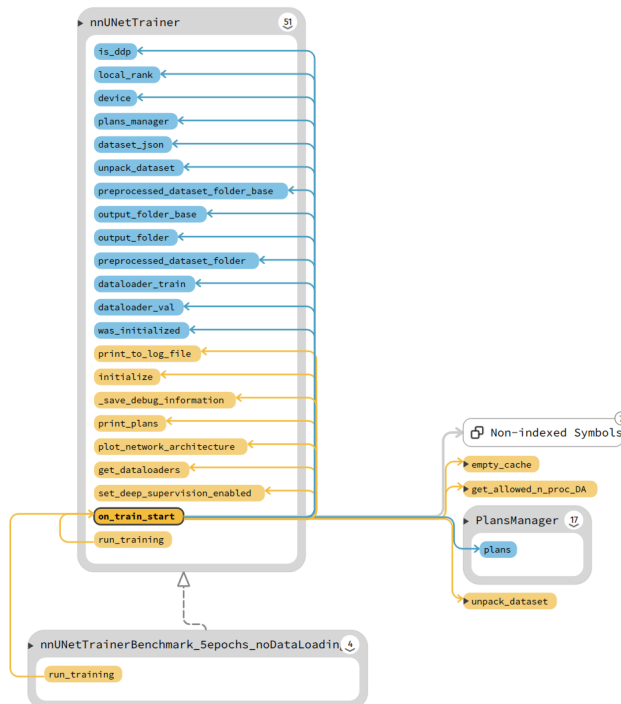
a. General idea

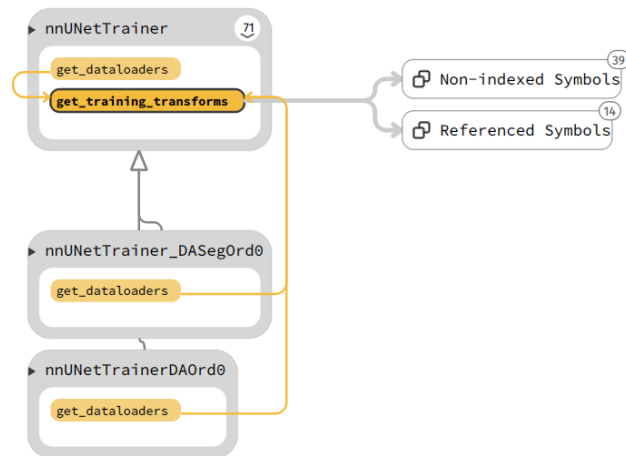
- As suggested by the author, data augmentation operations are mostly **fixed** (while nnUNet itself highlights adaptive data processing and training policies tailored to individual datasets).
- Based on commonly used **spatial transforms** and **color transforms**. Simple, conventional, but shown to be very effective in most cases (common anatomical structures and lesions + low-to-intermediate covariate shifts).
- Heavily relying on [batchgenerators](<https://github.com/MIC-DKFZ/batchgenerators>), a separate package written by DKFZ. Implemented mostly on CPU: numpy, scikit-image, etc. Further acceleration by their home-made multi-thread handlers (in practice people also use acceleration by pytorch dataloaders. while additional attention is needed when handling numpy random states (!) in that case).

b. Where and how it is called

- `/nnunetv2/training/nnUNetTrainer/nnUNetTrainer.py:`

`run_training` → `on_training_start` → `get_data loaders` → `get_training_transforms`





```

566 def get_dataloaders(self):
567     # we use the patch size to determine whether we need 2D or 3D dataloaders. We also use it to determine whether
568     # we need to use dummy 2D augmentation (in case of 3D training) and what our initial patch size should be
569     patch_size = self.configuration_manager.patch_size
570     dim = len(patch_size)
571
572     # needed for deep supervision: how much do we need to downscale the segmentation targets for the different
573     # outputs?
574     deep_supervision_scales = self._get_deep_supervision_scales()
575
576     rotation_for_DA, do_dummy_2d_data_aug, initial_patch_size, mirror_axes = \
577         self.configure_rotation_dummyDA_mirroring_and_initial_patch_size()
578
579     # training pipeline
580     tr_transforms = self.get_training_transforms(
581         patch_size, rotation_for_DA, deep_supervision_scales, mirror_axes, do_dummy_2d_data_aug,
582         order_resampling_data=3, order_resampling_seg=1,
583         use_mask_for_norm=self.configuration_manager.use_mask_for_norm,
584         is_cascaded=self.is_cascaded, foreground_labels=self.label_manager.foreground_labels,
585         regions=self.label_manager.foreground_regions if self.label_manager.has_regions else None,
586         ignore_label=self.label_manager.ignore_label)
587
588     # validation pipeline
589     val_transforms = self.get_validation_transforms(deep_supervision_scales,
590         is_cascaded=self.is_cascaded,
591         foreground_labels=self.label_manager.foreground_labels,
592         regions=self.label_manager.foreground_regions if
593             self.label_manager.has_regions else None,
594         ignore_label=self.label_manager.ignore_label)
595
596     dl_tr, dl_val = self.get_plain_dataloaders(initial_patch_size, dim)
597
598     allowed_num_processes = get_allowed_n_proc_DA()
599     if allowed_num_processes == 0:
600         mt_gen_train = SingleThreadedAugmenter(dl_tr, tr_transforms)
601         mt_gen_val = SingleThreadedAugmenter(dl_val, val_transforms)
602     else:
603         mt_gen_train = LimitedLenWrapper(self.num_iterations_per_epoch, data_loader=dl_tr, transform=tr_transforms,
604             num_processes=allowed_num_processes, num_cached=6, seeds=None,
605             pin_memory=self.device.type == 'cuda', wait_time=0.02)
606         mt_gen_val = LimitedLenWrapper(self.num_val_iterations_per_epoch, data_loader=dl_val,
607             transform=val_transforms, num_processes=max(1, allowed_num_processes // 2),
608             num_cached=3, seeds=None, pin_memory=self.device.type == 'cuda',
609             wait_time=0.02)
610
611     return mt_gen_train, mt_gen_val

```

Summary

- In `/nnunetv2/training/nnUNetTrainer/nnUNetTrainer.py`
 1. Line 576: Deciding rotation policy (adaptive) + computing actual patch size after possible rotations.
 2. Line 580:

Spatial transforms (almost fixed).

Intensity (color) transforms (almost fixed).

Additional steps for cascaded training (low-res + full_res).

Will be introduced in details in the below section.

2. Key components

a. Deciding rotation policies and possible patch sizes (obtained after rotations)

Where it resides:

`nnUNetTrainer.py`:

```
566 | def get_data_loaders(self):
567 |     # we use the patch size to determine whether we need 2D or 3D data loaders. We also use it to determine whether
568 |     # we need to use dummy 2D augmentation (in case of 3D training) and what our initial patch size should be
569 |     patch_size = self.configuration_manager.patch_size
570 |     dim = len(patch_size)
571 |
572 |     # needed for deep supervision: how much do we need to downscale the segmentation targets for the different
573 |     # outputs?
574 |     deep_supervision_scales = self._get_deep_supervision_scales()
575 |
576 |     rotation_for_DA, do_dummy_2d_data_aug, initial_patch_size, mirror_axes = \
577 |         self.configure_rotation_dummyDA_mirroring_and_inital_patch_size()
---
```

This calls the following:

```
566 | def get_data_loaders(self):
567 |     # we use the patch size to determine whether we need 2D or 3D data loaders. We also use it to determine whether
568 |     # we need to use dummy 2D augmentation (in case of 3D training) and what our initial patch size should be
569 |     patch_size = self.configuration_manager.patch_size
570 |     dim = len(patch_size)
571 |
572 |     # needed for deep supervision: how much do we need to downscale the segmentation targets for the different
573 |     # outputs?
574 |     deep_supervision_scales = self._get_deep_supervision_scales()
575 |
576 |     rotation_for_DA, do_dummy_2d_data_aug, initial_patch_size, mirror_axes = \
577 |         self.configure_rotation_dummyDA_mirroring_and_inital_patch_size()
---
```

The detailed implementation:

```

354 def configure_rotation_dummyDA_mirroring_and_initial_patch_size(self):
355     """
356     This function is stupid and certainly one of the weakest spots of this implementation. Not entirely sure how we can f
357     """
358     patch_size = self.configuration_manager.patch_size
359     dim = len(patch_size)
360     # todo rotation should be defined dynamically based on patch size (more isotropic patch sizes = more rotation)
361     if dim == 2:
362         do_dummy_2d_data_aug = False
363         # todo revisit this parametrization
364         if max(patch_size) / min(patch_size) > 1.5:
365             rotation_for_DA = {
366                 'x': (-15. / 360 * 2. * np.pi, 15. / 360 * 2. * np.pi),
367                 'y': (0, 0),
368                 'z': (0, 0)
369             }
370         else:
371             rotation_for_DA = {
372                 'x': (-180. / 360 * 2. * np.pi, 180. / 360 * 2. * np.pi),
373                 'y': (0, 0),
374                 'z': (0, 0)
375             }
376         mirror_axes = (0, 1)
377     elif dim == 3:
378         # todo this is not ideal. We could also have patch_size (64, 16, 128) in which case a full 180deg 2d rot would be
379         # order of the axes is determined by spacing, not image size
380         do_dummy_2d_data_aug = (max(patch_size) / patch_size[0]) > ANISO_THRESHOLD
381         if do_dummy_2d_data_aug:
382             # why do we rotate 180 deg here all the time? We should also restrict it
383             rotation_for_DA = {
384                 'x': (-180. / 360 * 2. * np.pi, 180. / 360 * 2. * np.pi),
385                 'y': (0, 0),
386                 'z': (0, 0)
387             }
388         else:
389             rotation_for_DA = {
390                 'x': (-30. / 360 * 2. * np.pi, 30. / 360 * 2. * np.pi),
391                 'y': (-30. / 360 * 2. * np.pi, 30. / 360 * 2. * np.pi),
392                 'z': (-30. / 360 * 2. * np.pi, 30. / 360 * 2. * np.pi),
393             }
394         mirror_axes = (0, 1, 2)
395     else:
396         raise RuntimeError()
397
398     # todo this function is stupid. It doesn't even use the correct scale range (we keep things as they were in the
399     # old nnunet for now)
400     initial_patch_size = get_patch_size(patch_size[-dim:],
401                                         *rotation_for_DA.values(),
402                                         (0.85, 1.25))
403     if do_dummy_2d_data_aug:
404         initial_patch_size[0] = patch_size[0]
405
406     self.print_to_log_file(f'do_dummy_2d_data_aug: {do_dummy_2d_data_aug}')
407     self.inference_allowed_mirroring_axes = mirror_axes
408
409     return rotation_for_DA, do_dummy_2d_data_aug, initial_patch_size, mirror_axes

```

Summary:

- `/nnunetv2/training/nnUNetTrainer/nnUNetTrainer.py` Line 358 - Line 399:

Deciding the rotation policy based on the **anisotropy** (how patch size along different dimensions may differ) of pre-configured patch size along different dimensions.

This is controlled by `do_dummy_2d_data_aug`: if the patch size has high anisotropy, then restrict the range of the rotation angle:

For example, a 90-degree rotation for a rectangular patch with a high aspect ratio may not make too much sense: much of the expanded patch after rotation will be empty. Therefore their range of rotation along the dimensions with high anisotropy would be restrained to +- 15 degree for 2D or +- 30 degrees for 3D.

(As suggested by the author, order of the axes in `patch_size` is determined by spacing: smaller ones first.)

- `/nnunetv2/training/nnUNetTrainer/nnUNetTrainer.py` Line 400 - 403:

Computing the **largest actual patch size (final_shape)** you can get after possible rotations.

This calls `/nnunetv2/training/data_augmentation/compute_initial_patch_size.py`:

```
4 def get_patch_size(final_patch_size, rot_x, rot_y, rot_z, scale_range):
5     if isinstance(rot_x, (tuple, list)):
6         rot_x = max(np.abs(rot_x))
7     if isinstance(rot_y, (tuple, list)):
8         rot_y = max(np.abs(rot_y))
9     if isinstance(rot_z, (tuple, list)):
10        rot_z = max(np.abs(rot_z))
11    rot_x = min(90 / 360 * 2. * np.pi, rot_x)
12    rot_y = min(90 / 360 * 2. * np.pi, rot_y)
13    rot_z = min(90 / 360 * 2. * np.pi, rot_z)
14    from batchgenerators.augmentations.utils import rotate_coords_3d, rotate_coords_2d
15    coords = np.array(final_patch_size)
16    final_shape = np.copy(coords)
17    if len(coords) == 3:
18        final_shape = np.max(np.vstack((np.abs(rotate_coords_3d(coords, rot_x, 0, 0)), final_shape)), 0)
19        final_shape = np.max(np.vstack((np.abs(rotate_coords_3d(coords, 0, rot_y, 0)), final_shape)), 0)
20        final_shape = np.max(np.vstack((np.abs(rotate_coords_3d(coords, 0, 0, rot_z)), final_shape)), 0)
21    elif len(coords) == 2:
22        final_shape = np.max(np.vstack((np.abs(rotate_coords_2d(coords, rot_x)), final_shape)), 0)
23    final_shape /= min(scale_range)
24    return final_shape.astype(int)
```

- Line 660 & 680: compress the channel dimension if only b,c, x, y, z → b,c*x,y,z

b. Spatial transforms

Overview

Most of transforms are implemented in [batchgenerators](<https://github.com/MIC-DKFZ/batchgenerators>).

In batchgenerators, callables are defined in `batchgenerators/batchgenerators/transforms/`

Actual implementations of transform functions are in `batchgenerators/batchgenerators/augmentations/`

Mostly rigid transforms. Elastic transforms are `disabled` by default (why?). Need to manually enable them. In `nnUnetv1`, DA configurations are detached from the trainer so you can directly change them in the configuration file to accommodate your own changes. In `v2` you probably need to derive a variant trainer class (as those in `nnunetv2/training/nnUNetTrainer/variants/`) to do so.

Where it resides

- `/nnunetv2/training/nnUNetTrainer/nnUNetTrainer.py`

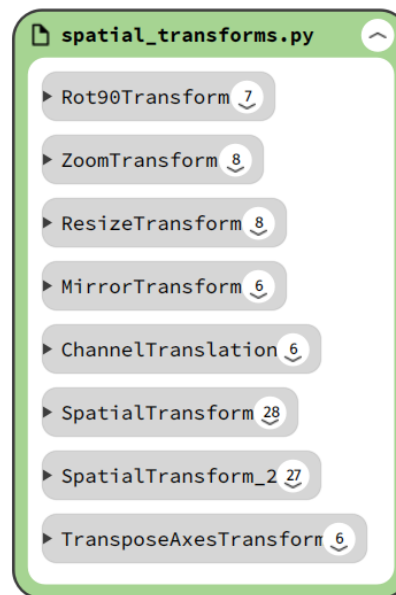
```

666     tr_transforms.append(SpatialTransform(
667         patch_size_spatial, patch_center_dist_from_border=None,
668         do_elastic_deform=False, alpha=(0, 0), sigma=(0, 0),
669         do_rotation=True, angle_x=rotation_for_DA['x'], angle_y=rotation_for_DA['y'], angle_z=rotation_for_DA['z'],
670         p_rot_per_axis=1, # todo experiment with this
671         do_scale=True, scale=(0.7, 1.4),
672         border_mode_data="constant", border_cval_data=0, order_data=order_resampling_data,
673         border_mode_seg="constant", border_cval_seg=border_val_seg, order_seg=order_resampling_seg,
674         random_crop=False, # random cropping is part of our dataloaders
675         p_el_per_sample=0, p_scale_per_sample=0.2, p_rot_per_sample=0.2,
676         independent_scale_for_each_axis=False # todo experiment with this
677     ))
678
679     if do_dummy_2d_data_aug:
680         tr_transforms.append(Convert2DTo3DTransform())

```

This calls augmentation functions in batchgenerators

`/batchgenerators/transforms/spatial_transforms.py` → `/batchgenerators/augmentations/spatial_transformations.py`



Key facts:

- Use `data_key` and `label_key` to route images or labels to different operations (e.g. the order of interpolation. When you override their implementation, just keep the interface the same.)

```

249 | class SpatialTransform(AbstractTransform):
    |
    |
332 |     def __call__(self, **data_dict):
333 |         data = data_dict.get(self.data_key)
334 |         seg = data_dict.get(self.label_key)
335 |

```

- No elastic transforms by default.
- Handling with padded borders: In most cases for the default pre-processing configuration, the background/boundary of pre-processed images are set to 0. Therefore the padded values are also 0 (Line

672, also the original paper).

Padding values for segmentation masks are set to -1 (Line 673) and later changed back to 0 (Line 701).

- Resampling order = 3 for images, 0 (nearest-neighbor) for segmentation labels.

c. Intensity/color augmentations

Overview

These augmentation functions are also based on [batchgenerators], with low-level implementations on CPU (numpy, scikit-images, etc.)

Most of the augmentation functions are commonly used. A side reference is [bigaug] (<https://ieeexplore.ieee.org/document/8995481>) which experiments the effects of individual augmentation operations on segmentation performances (incl. those with out-of-domain test samples).

Where it resides

- `/nnunetv2/training/nnUNetTrainer/nnUNetTrainer.py`

```
682     tr_transforms.append(GaussianNoiseTransform(p_per_sample=0.1))
683     tr_transforms.append(GaussianBlurTransform((0.5, 1.), different_sigma_per_channel=True, p_per_sample=0.2,
684                                             p_per_channel=0.5))
685     tr_transforms.append(BrightnessMultiplicativeTransform(multiplier_range=(0.75, 1.25), p_per_sample=0.15))
686     tr_transforms.append(ContrastAugmentationTransform(p_per_sample=0.15))
687     tr_transforms.append(SimulateLowResolutionTransform(zoom_range=(0.5, 1), per_channel=True,
688                                                         p_per_channel=0.5,
689                                                         order_downsample=0, order_upsample=3, p_per_sample=0.25,
690                                                         ignore_axes=ignore_axes))
691     tr_transforms.append(GammaTransform((0.7, 1.5), True, True, retain_stats=True, p_per_sample=0.1))
692     tr_transforms.append(GammaTransform((0.7, 1.5), False, True, retain_stats=True, p_per_sample=0.3))
693
```

These augmentation callables are defined in `/batchgenerators/transforms/color_transforms.py` and `/batchgenerators/transforms/noise_transforms.py`

Interesting facts

- Gamma transforms: applied `twice` (line 691-692), one on original intensities, one on inverted intensities (see `batchgenerators/augmentations/color_augmentations.py` line 109-110)

```
107 def augment_gamma(data_sample, gamma_range=(0.5, 2), invert_image=False, epsilon=1e-7, per_channel=False,
108                  retain_stats: Union[bool, Callable[[], bool]] = False):
109     if invert_image:
110         data_sample = - data_sample
```





Illustration of Gamma transform functions with different γ 's. Source: https://www.researchgate.net/figure/Gamma-transformation-of-grey-level_fig5_271632132

d. Additional augmentations for cascaded training (low-res)

Overview

Cascaded training:

“... a 3D U-Net cascade where first a 3D U-Net operates on **low resolution** images and then a second high-resolution 3D U-Net **refined** the predictions of the former (for 3D datasets with large image sizes only)...”

The pitfall for the cascaded scheme is that (based on the understanding of my own): if the low-res model yields perfect fitting in the training stage, the high-res U-Net may degrade and learns **identity mapping** instead. As a result, in testing, when the prediction of low-res model is not ideal, the high-res model may not be sufficiently robust to make corrections. The author therefore introduces stochasticity to the ground truth for the low-res model.

Another way of interpreting: there may be multiple possible high-res segmentations corresponding to the same low-res image.

To this end, the author makes perturbations: dilation, erosion, opening, closing, and noises to the low-resolution segmentation labels.

Where it resides

- `/nnunetv2/training/nnUNetTrainer/nnUNetTrainer.py`

```

703         if is_cascaded:
704             assert foreground_labels is not None, 'We need foreground_labels for cascade augmentations'
705             tr_transforms.append(MoveSegAsOneHotToData(1, foreground_labels, 'seg', 'data'))
706             tr_transforms.append(ApplyRandomBinaryOperatorTransform(
707                 channel_idx=list(range(-len(foreground_labels), 0)),
708                 p_per_sample=0.4,
709                 key="data",
710                 strel_size=(1, 8),
711                 p_per_label=1))
712             tr_transforms.append(
713                 RemoveRandomConnectedComponentFromOneHotEncodingTransform(
714                     channel_idx=list(range(-len(foreground_labels), 0)),
715                     key="data",
716                     p_per_sample=0.2,
717                     fill_with_other_class_p=0,
718                     dont_do_if_covers_more_than_x_percent=0.15))

```

Key facts

- `MoveSegAsOneHotToData`: get the one_hot labels of the foreground.
- `ApplyRandomBinaryOperatorTransform`: perform random perturbations: dilation, erosion, closing, opening... to one-hot labels.

- `RemoveRandomConnectedComponentFromOne ...`

Randomly removing connected components with size smaller than a certain range and/or move that to other classes.

Connected component operations come from [acvl_utils](https://github.com/MIC-DKFZ/acvl_utils/blob/master/acvl_utils/morphology/morphology_helper.py)

Why? Simulating erroneous predictions for improving robustness (?my own guess).

- Additional steps for region-based training (merging multiple labels to one region, specific to BraTS):

```

722         if regions is not None:
723             # the ignore label must also be converted
724             tr_transforms.append(ConvertSegmentationToRegionsTransform(list(regions) + [ignore_label]
725                 if ignore_label is not None else regions,
726                 'target', 'target'))

```

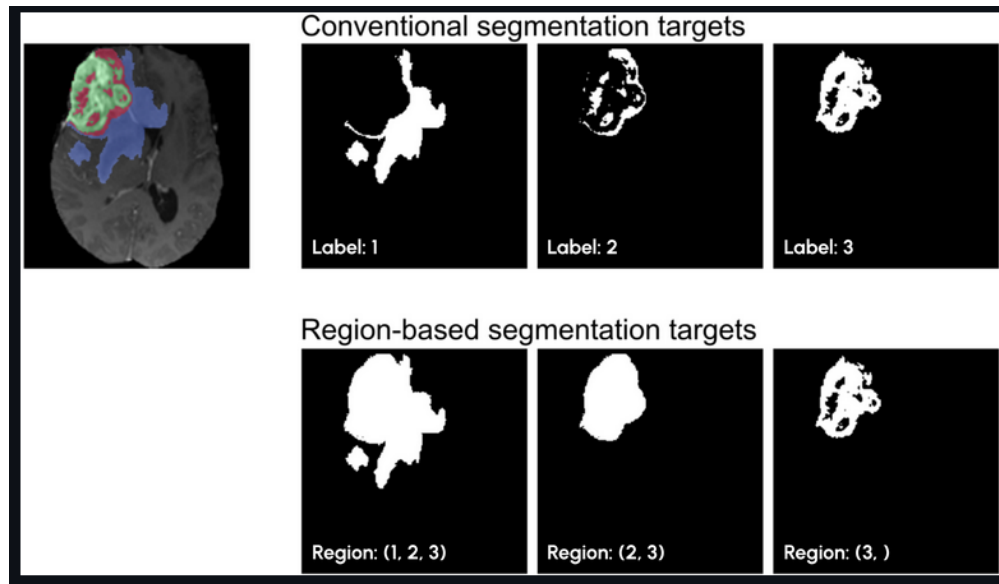


Illustration of region-based-training, source: https://github.com/MIC-DKFZ/nnUNet/blob/master/documentation/region_based_training.md

- Lower resolutions for deep supervision (in practice, very helpful for domain robustness). Implemented using `scikit-image`.

```
728 |         if deep_supervision_scales is not None:
729 |             tr_transforms.append(DownsampleSegForDSTransform2(deep_supervision_scales, 0, input_key='target',
730 |                                                                 output_key='target'))
```

- Output format: Now goes from numpy array to pytorch tensor.

```
731 |         tr_transforms.append(NumpyToTensor(['data', 'target'], 'float'))
```

3. Comments

Advantages:

- Off-the-shelf augmentation configurations that work well for most cases, `little domain-specific knowledge` is needed for the dataset.

Not sure how the developers come up with this set of configurations (It is a large search space).

- Suitable for most of `medical image segmentation challenges`: giving you a good-enough baseline in the shortest time without having you wasting additional efforts on tuning data augmentation configurations.

Comments:

- If stronger augmentations are needed (based on **adequate understanding of your datasets and tasks**), variants can be found in `nnunetv2/training/nnUNetTrainer/variants/data_augmentation/` (`moreDA` and `insaneDA` for v1 code)
(e.g. domain generalization-related challenges favors stronger/task-specific DA)
- Two levels of parallelization: home-made ones and those by native `numpy`. Be aware of that when sharing CPU resources with other users.
- The augmentation part `alone` may not always give you the same magic compared with nnUNet as a whole. (Same conclusion from the author of nnUNet).

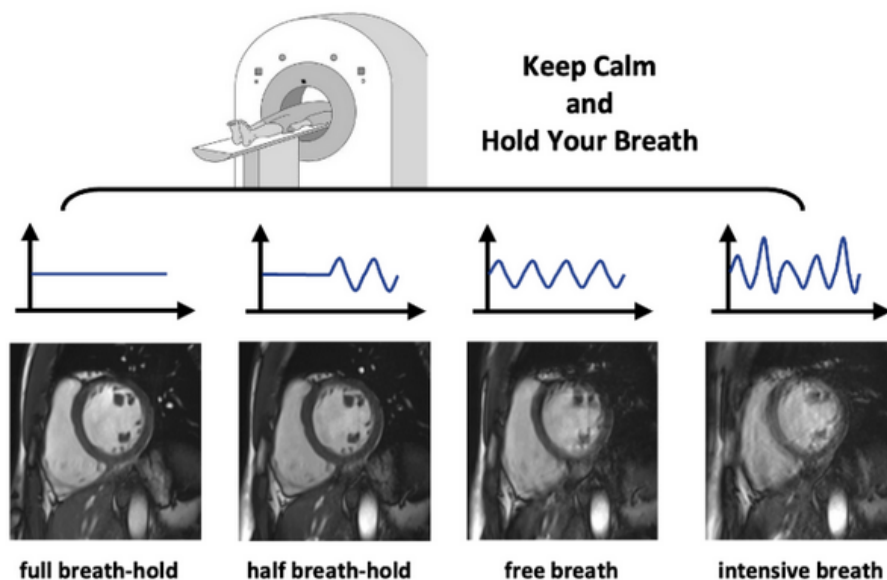
4. Beyond pre-configured data augmentations

Data augmentation is crucial for tasks that relate to `low-resource training`, `semi-supervised learning`, and `domain generalization`. In those cases, more advanced data augmentations can be tailored to the specific tasks.

Taking CMRxMotion and FeTA 2022 as examples. The winning solutions to them do employ nnUNet **but** importantly, they tailor data augmentation to their specific datasets and tasks.

The core idea for designing augmentation policy is to `simulate the diversity of the real underlying data`, based on generic assumptions about medical images and/or on your domain knowledge of the task at hand:

- [CMR-Motion](<http://cmr.miccai.cloud/>):

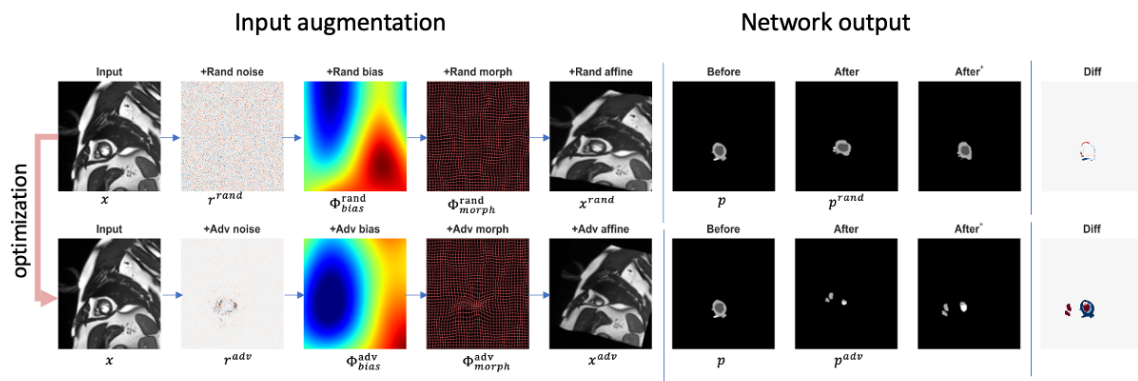


Task: Segmenting short axis CINE cardiac MRIs that may contain unexpected respiratory artifacts (see the right three examples). Essentially a **domain generalization** challenge.

Analysis: The performance bottleneck is to handle **unforeseen respiratory** in MRI. Beyond these motion patterns, other stochastics caused by imaging acquisition processes exist, and lead to a diverse test dataset that may be different from the training data.


Winning solution (CUHK, two components):

1. augmentation techniques that **simulates the stochastics in the physical MRI acquisition process** [AdvChain](<https://github.com/cherise215/advchain>): bias field, noises, elastic deformations, etc. **Adversarial training** is used to optimize the augmentation parameters of data augmentation operations.




2. Pre-training on external datasets (ACDC, MSCMRSeg, etc), as allowed by the organizers

- [FeTA](<https://feta.grand-challenge.org/>)




25th International Conference on
Medical Image Computing and
Computer Assisted Intervention
September 18-22, 2022
Resorts World Convention Centre Singapore

Challenges in Fetal Brain Segmentation



- Different institutions
 - Different MR scanners
 - Different image acquisition protocols
 - Different post-processing tools
 - Super-Resolution Reconstructions
 - Different image resolutions



Domain Generalization Problem

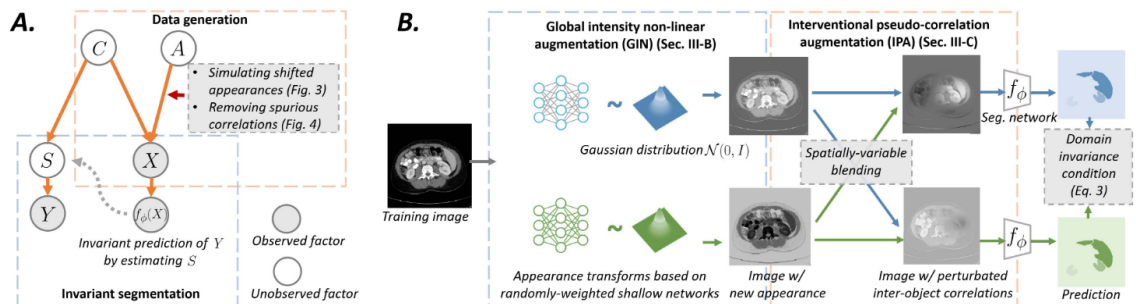
Task: To segment fetal brain MRIs that may come from **unforeseen** institutes. Essentially also a **domain generalization** challenge.

Analysis: The performance bottleneck is to handle **unforeseen intensity distributions**.

Winning solution (TUM & ICL): employing augmentation techniques that cover common distribution shifts in image intensities across datasets, these augmentation includes:

1. Conventional brightness/contrast/random Gamma, etc.
2. Bias-field (specific to MRI, same as described in [AdvChain])
3. Generic and strong domain generalization technique:

[Causal-SD](<https://github.com/cheng-01037/Causality-Medical-Image-Domain-Generalization>)



They also employ **ensembling strategies** among model trained with different augmentation policies, for better overall performance for both in-distribution data and various types of out-of-distribution data (as we do not know whether a test image is in-distribution or out-of-distribution, nor do we know the source of OOD).

Some of the ensemble rules are similar to those described in [Causal-SD].